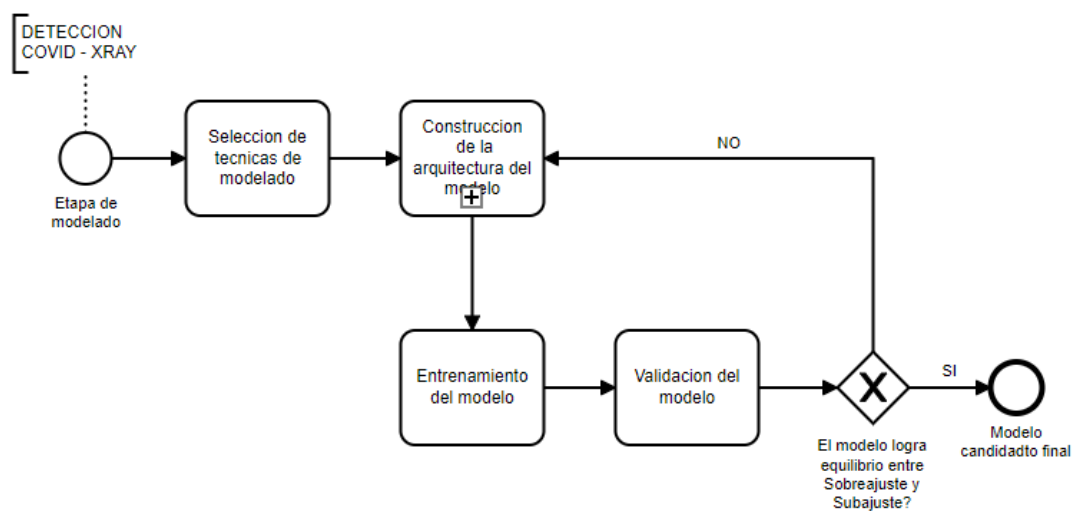


Modelado usando Transfer Learning con DenseNet169 y balanceo por Perdida Focal

August 29, 2021

Esta fase de la metodología consiste en extraer el valor de los datos desarrollando un modelo que aprenda de los patrones en estos.

- El diagrama en cuestion de esta fase esta a continuacion:



```
[1]: import tensorflow as tf
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from tensorflow.keras.preprocessing.image import ImageDataGenerator #generator_
    ↳ de imagenes
from sklearn.utils import class_weight
from sklearn.metrics import classification_report
import random
```

```
[2]: tf.keras.backend.clear_session()
```

```
[3]: tf.random.set_seed(42)#semilla 42, para la reproducibilidad de resultados
      random.seed(42)
      np.random.seed(42)
```

0.0.1 Canalizacion de datos

- Preparamos la canalizacion de datos, a partir de las imagenes del disco.

```
[4]: train_datagen=ImageDataGenerator(
      rescale=1.0/255, #escalamos los datos en rangos de [-1,1] El modelo
      ↪ MobileNetv2 espera esta configuracion
      rotation_range = 45,
      zoom_range = 0.2,
      shear_range = 0.2,
      width_shift_range = 0.2,
      height_shift_range = 0.2,
      horizontal_flip=True,
      vertical_flip = True,
      fill_mode = 'nearest'
    )
#sobre los datos de validacion y test no se hace ningun aumento de datos.
    validation_datagen=ImageDataGenerator(rescale=1.0/255) #escalamiento de
    ↪ validacion a un rango de [0,1]
    test_datagen=ImageDataGenerator(rescale=1.0/255) #escalamiento de test a
    ↪ un rango de [0,1]
```

```
[5]: #definimos las rutas para el acceso a los datos
      train_path="../input/datasetv3/Datasets/train"
      validation_path="../input/datasetv3/Datasets/val"
      test_path="../input/datasetv3/Datasets/test"

      #creamos los generadores de datos a partir de los flujos de informacion
      BATCH_SIZE=32 #tamaño del lote que se ira pasando poco a poco
      IMAGE_SIZE=(256,256)

      train_generator=train_datagen.flow_from_directory(
          train_path,
          target_size=IMAGE_SIZE,
          batch_size=BATCH_SIZE,
          class_mode="categorical"
      )

      validation_generator=validation_datagen.flow_from_directory(
          validation_path,
          target_size=IMAGE_SIZE,
          batch_size=BATCH_SIZE,
          class_mode="categorical"
```

```

)

test_generator=test_datagen.flow_from_directory(
    test_path,
    target_size=IMAGE_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="categorical"
)

```

Found 15238 images belonging to 3 classes.

Found 1694 images belonging to 3 classes.

Found 4233 images belonging to 4 classes.

0.0.2 Técnica para el tratamiento de datos Desbalanceados. Focal Loss Categorical

- Para ello personalizaremos una función de costo
- En nuestro caso trabajamos con imágenes, y tenemos la opción de tratar este desbalanceo penalizando los pesos de las clases mayoritarias a favor de las clases minoritarias.

```

[6]: def focal_loss(gamma=2., alpha=4.):

    gamma = float(gamma)
    alpha = float(alpha)

    def focal_loss_fixed(y_true, y_pred):
        """Focal loss for multi-classification
        FL(p_t)=-alpha(1-p_t)^(gamma)ln(p_t)
        Notice: y_pred is probability after softmax
        gradient is d(FL)/d(p_t) not d(FL)/d(x) as described in paper
        d(FL)/d(p_t) * [p_t(1-p_t)] = d(FL)/d(x)
        Focal Loss for Dense Object Detection
        https://arxiv.org/abs/1708.02002
        Arguments:
            y_true {tensor} -- ground truth labels, shape of [batch_size,
↪num_cls]
            y_pred {tensor} -- model's output, shape of [batch_size, num_cls]
        Keyword Arguments:
            gamma {float} -- (default: {2.0})
            alpha {float} -- (default: {4.0})
        Returns:
            [tensor] -- loss.
        """
        epsilon = 1.e-9
        y_true = tf.convert_to_tensor(y_true, tf.float32)
        y_pred = tf.convert_to_tensor(y_pred, tf.float32)

        model_out = tf.add(y_pred, epsilon)
        ce = tf.multiply(y_true, -tf.math.log(model_out))

```

```

weight = tf.multiply(y_true, tf.pow(tf.subtract(1., model_out), gamma))
fl = tf.multiply(alpha, tf.multiply(weight, ce))
reduced_fl = tf.reduce_max(fl, axis=1)
return tf.reduce_mean(reduced_fl)
return focal_loss_fixed

```

0.0.3 Selección de tecnicas de Modelado

Al tratarse de un problema de clasificacion de imagenes entre los posibles candidatos tenemos:

- MultiLayer Perceptron: Red neuronal de capas densamente conectadas
- Convolutional Neural Network: Red neuronal convolucional.
- Modelos de machine learning clasico (Maquinas de soporte vectorial, arboles de decision e impulso, etc.)

Escogi la red neuronal convolucional porque **aprende de patrones locales** como rasgos pequeños y en bloques de informacion, mientras que el **MLP** aprende de patrones especificos, e decir de todo el espacio de entrada en general.

0.0.4 Construcción de la arquitectura del modelo

- Para la construccion de la arquitectura crearemos un modelo desde 0 con una arquitectura de red neuronal solida.

Una vez obtenidos la arquitectura del modelo * Generando **Callbacks** para detener el entrenamiento cuando no se tienen buenos resultados

```

[7]: #cuando la funcion de perdida ya no mejora en los datos de validacion.
early_stopping=tf.keras.callbacks.EarlyStopping(monitor="val_loss",patience=10,
                                                mode="min",
                                                restore_best_weights=True)
model_checkpoint=tf.keras.callbacks.ModelCheckpoint("base_model_best_check.
↪h5",save_best_only=True)
n_epochs=100

```

1 Aplicación de Transfer Learning usando DenseNet169 y tecnica de anti-desbalanceo Focal Loss Categorical

- Se escogio la arquitectura DenseNet169 por tener un mayor precision sobre el conjunto de datos Image.net en la que fue entrenado.

```

[8]: INPUT_SHAPE=(256,256,3)
BASE_LEARNING_RATE=0.0001
def build_model_transferLearning():
    dense_net169=tf.keras.applications.DenseNet169(
        weights="imagenet",
        input_shape=INPUT_SHAPE,
        include_top=False
    )

```

```

#descongelamos algunas capas
dense_net169.trainable=True
for layer in dense_net169.layers:
    if 'conv5' in layer.name:
        layer.trainable=True
    else:
        layer.trainable=False
#creamos el modelo
inputs=tf.keras.Input(shape=INPUT_SHAPE)
x=dense_net169(inputs)
x=tf.keras.layers.Flatten()(x)

x=tf.keras.layers.BatchNormalization()(x)
x=tf.keras.layers.Dense(256,activation="relu")(x)
x=tf.keras.layers.Dropout(0.4)(x)

x=tf.keras.layers.BatchNormalization()(x)
x=tf.keras.layers.Dense(128,activation="relu")(x)
x=tf.keras.layers.Dropout(0.4)(x)

outputs=tf.keras.layers.Dense(3,activation="softmax")(x)

model=tf.keras.Model(inputs,outputs)

model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=BASE_LEARNING_RATE),
    loss=focal_loss(alpha=1.0),
    metrics=[
        tf.keras.metrics.CategoricalAccuracy(name="accuracy"),
        tf.keras.metrics.Recall(name="recall")
    ]
)
return model

```

```
[9]: transfer_model=build_model_transferLearning()
```

```
[10]: transfer_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 256, 256, 3)]	0
densenet169 (Functional)	(None, 8, 8, 1664)	12642880
flatten (Flatten)	(None, 106496)	0

batch_normalization (BatchNo	(None, 106496)	425984

dense (Dense)	(None, 256)	27263232

dropout (Dropout)	(None, 256)	0

batch_normalization_1 (Batch	(None, 256)	1024

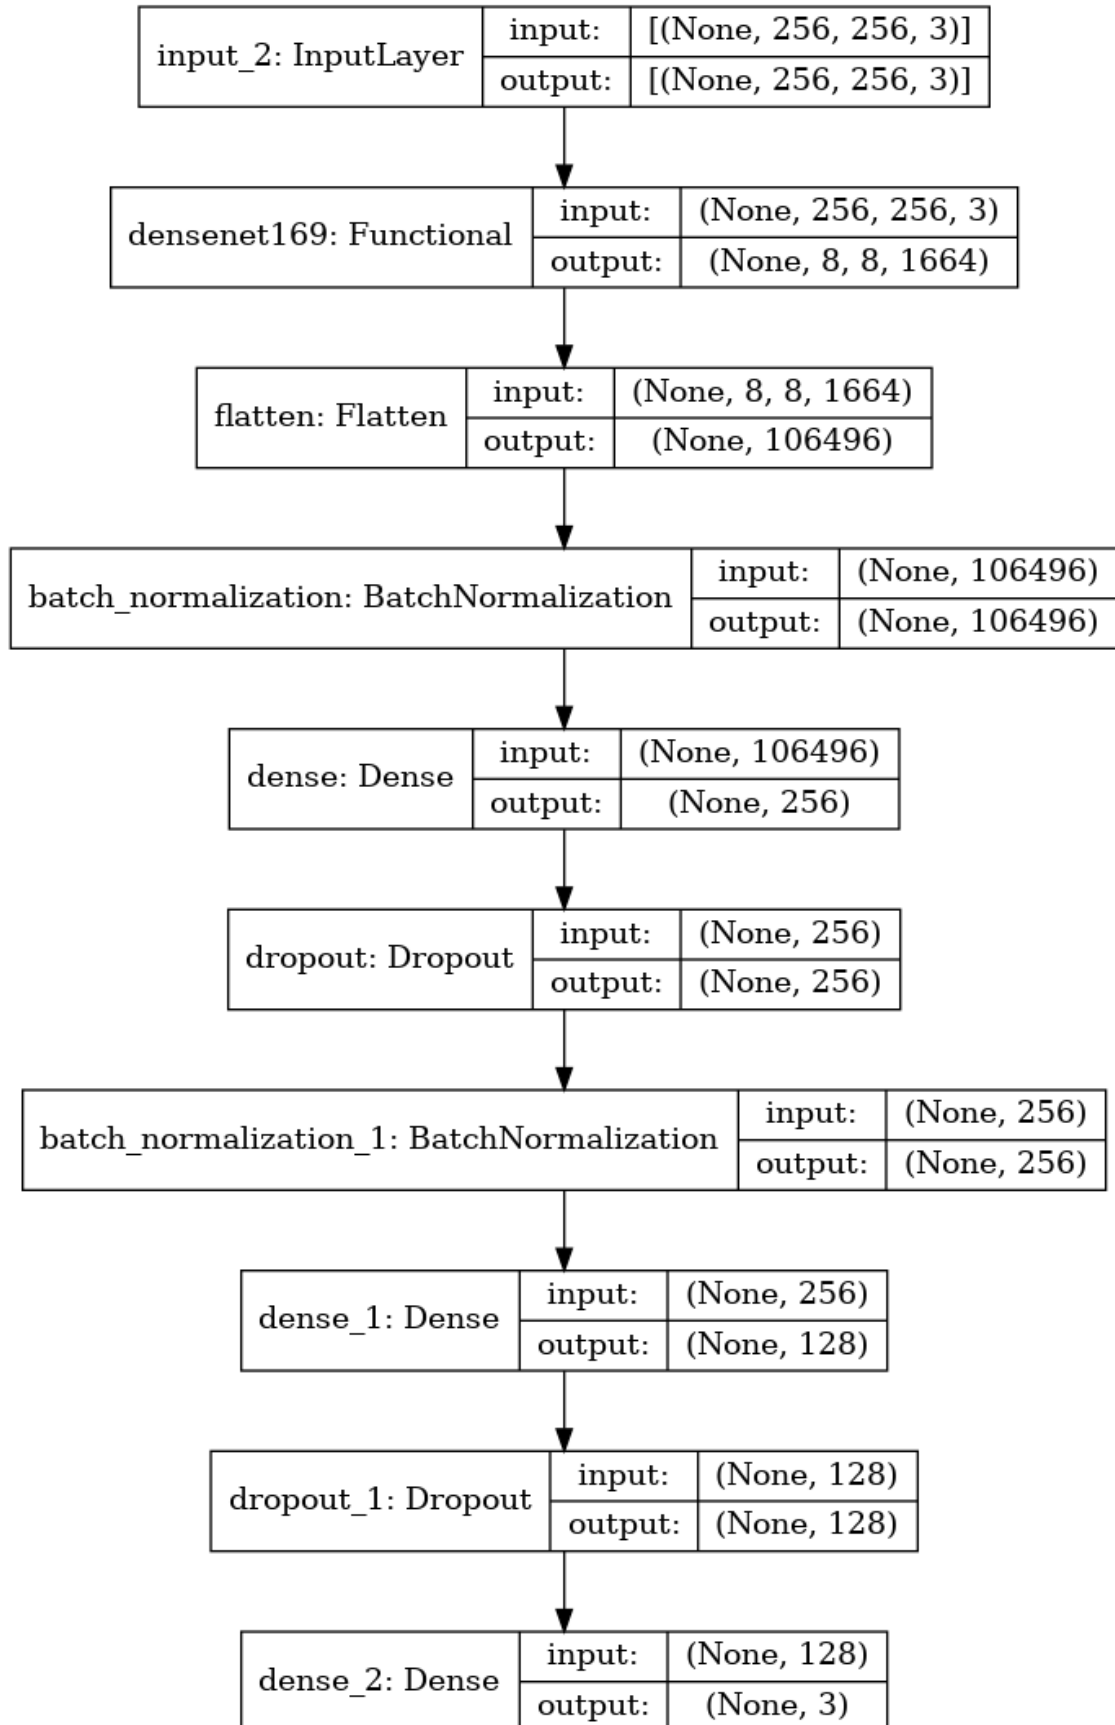
dense_1 (Dense)	(None, 128)	32896

dropout_1 (Dropout)	(None, 128)	0

dense_2 (Dense)	(None, 3)	387
=====		
Total params: 40,366,403		
Trainable params: 33,423,619		
Non-trainable params: 6,942,784		

```
[11]: plot_model(transfer_model, "transfer_densenet.png", show_shapes=True)
```

```
[11]:
```



ENTRENAMIENTO DEL MODELO

- El entrenamiento se realiza en **100 épocas**, un generador de datos de entrenamiento, un generador de datos de validación, 2 callbacks para detener el entrenamiento de manera temprana en caso no se obtengan buenos resultados en base a la función de pérdida en los datos de validación durante 10 épocas consecutivas, y otra para guardar por puntos el mejor modelo obtenido hasta el momento.

```
[12]: transfer_model=build_model_transferLearning()
      history_model=transfer_model.fit(
          train_generator,
          epochs=n_epochs,
          validation_data=validation_generator,
          validation_steps=validation_generator.samples//
      ↪ validation_generator.batch_size,
          callbacks=[early_stopping,model_checkpoint],
          #class_weight=class_weights,#penalización de pesos para el
      ↪ balanceo,
          workers=8
      )
```

Epoch 1/100

477/477 [=====] - 252s 492ms/step - loss: 0.5307 - accuracy: 0.6529 - recall: 0.6065 - val_loss: 0.1964 - val_accuracy: 0.8540 - val_recall: 0.8305

Epoch 2/100

477/477 [=====] - 236s 490ms/step - loss: 0.2916 - accuracy: 0.7745 - recall: 0.7330 - val_loss: 0.1509 - val_accuracy: 0.8438 - val_recall: 0.7704

Epoch 3/100

477/477 [=====] - 234s 484ms/step - loss: 0.2221 - accuracy: 0.8032 - recall: 0.7576 - val_loss: 0.1360 - val_accuracy: 0.8630 - val_recall: 0.8269

Epoch 4/100

477/477 [=====] - 230s 477ms/step - loss: 0.1927 - accuracy: 0.8192 - recall: 0.7781 - val_loss: 0.1048 - val_accuracy: 0.8930 - val_recall: 0.8564

Epoch 5/100

477/477 [=====] - 233s 482ms/step - loss: 0.1743 - accuracy: 0.8372 - recall: 0.7931 - val_loss: 0.1109 - val_accuracy: 0.8756 - val_recall: 0.8287

Epoch 6/100

477/477 [=====] - 233s 483ms/step - loss: 0.1487 - accuracy: 0.8484 - recall: 0.8061 - val_loss: 0.1421 - val_accuracy: 0.8299 - val_recall: 0.7855

Epoch 7/100
477/477 [=====] - 236s 488ms/step - loss: 0.1386 -
accuracy: 0.8592 - recall: 0.8193 - val_loss: 0.0940 - val_accuracy: 0.8978 -
val_recall: 0.8756
Epoch 8/100
477/477 [=====] - 232s 480ms/step - loss: 0.1263 -
accuracy: 0.8562 - recall: 0.8211 - val_loss: 0.1051 - val_accuracy: 0.8600 -
val_recall: 0.8347
Epoch 9/100
477/477 [=====] - 231s 479ms/step - loss: 0.1277 -
accuracy: 0.8674 - recall: 0.8354 - val_loss: 0.0973 - val_accuracy: 0.8816 -
val_recall: 0.8534
Epoch 10/100
477/477 [=====] - 232s 480ms/step - loss: 0.1201 -
accuracy: 0.8688 - recall: 0.8392 - val_loss: 0.1123 - val_accuracy: 0.8804 -
val_recall: 0.8576
Epoch 11/100
477/477 [=====] - 231s 478ms/step - loss: 0.1144 -
accuracy: 0.8750 - recall: 0.8404 - val_loss: 0.1042 - val_accuracy: 0.8900 -
val_recall: 0.8666
Epoch 12/100
477/477 [=====] - 232s 480ms/step - loss: 0.1144 -
accuracy: 0.8759 - recall: 0.8491 - val_loss: 0.0943 - val_accuracy: 0.8936 -
val_recall: 0.8546
Epoch 13/100
477/477 [=====] - 233s 482ms/step - loss: 0.1117 -
accuracy: 0.8834 - recall: 0.8607 - val_loss: 0.0912 - val_accuracy: 0.9069 -
val_recall: 0.8720
Epoch 14/100
477/477 [=====] - 232s 480ms/step - loss: 0.0990 -
accuracy: 0.8944 - recall: 0.8730 - val_loss: 0.0889 - val_accuracy: 0.9075 -
val_recall: 0.8846
Epoch 15/100
477/477 [=====] - 230s 477ms/step - loss: 0.0952 -
accuracy: 0.8912 - recall: 0.8695 - val_loss: 0.0923 - val_accuracy: 0.8978 -
val_recall: 0.8510
Epoch 16/100
477/477 [=====] - 229s 474ms/step - loss: 0.0966 -
accuracy: 0.8957 - recall: 0.8760 - val_loss: 0.1073 - val_accuracy: 0.8732 -
val_recall: 0.8480
Epoch 17/100
477/477 [=====] - 230s 475ms/step - loss: 0.0854 -
accuracy: 0.9011 - recall: 0.8843 - val_loss: 0.0856 - val_accuracy: 0.9069 -
val_recall: 0.8852
Epoch 18/100
477/477 [=====] - 230s 475ms/step - loss: 0.0858 -
accuracy: 0.9053 - recall: 0.8852 - val_loss: 0.0945 - val_accuracy: 0.9056 -
val_recall: 0.8816

Epoch 19/100
477/477 [=====] - 229s 475ms/step - loss: 0.0920 -
accuracy: 0.8984 - recall: 0.8776 - val_loss: 0.0773 - val_accuracy: 0.9105 -
val_recall: 0.8954
Epoch 20/100
477/477 [=====] - 230s 477ms/step - loss: 0.0832 -
accuracy: 0.9005 - recall: 0.8843 - val_loss: 0.1113 - val_accuracy: 0.8948 -
val_recall: 0.8828
Epoch 21/100
477/477 [=====] - 231s 478ms/step - loss: 0.0866 -
accuracy: 0.9001 - recall: 0.8840 - val_loss: 0.0743 - val_accuracy: 0.9249 -
val_recall: 0.9087
Epoch 22/100
477/477 [=====] - 229s 475ms/step - loss: 0.0786 -
accuracy: 0.9131 - recall: 0.8978 - val_loss: 0.0780 - val_accuracy: 0.9189 -
val_recall: 0.9093
Epoch 23/100
477/477 [=====] - 229s 475ms/step - loss: 0.0795 -
accuracy: 0.9101 - recall: 0.8961 - val_loss: 0.0961 - val_accuracy: 0.9105 -
val_recall: 0.8924
Epoch 24/100
477/477 [=====] - 228s 473ms/step - loss: 0.0757 -
accuracy: 0.9133 - recall: 0.9005 - val_loss: 0.0878 - val_accuracy: 0.9032 -
val_recall: 0.8948
Epoch 25/100
477/477 [=====] - 229s 473ms/step - loss: 0.0803 -
accuracy: 0.9152 - recall: 0.9022 - val_loss: 0.1007 - val_accuracy: 0.8924 -
val_recall: 0.8786
Epoch 26/100
477/477 [=====] - 229s 473ms/step - loss: 0.0731 -
accuracy: 0.9151 - recall: 0.9041 - val_loss: 0.0729 - val_accuracy: 0.9159 -
val_recall: 0.9032
Epoch 27/100
477/477 [=====] - 230s 475ms/step - loss: 0.0787 -
accuracy: 0.9152 - recall: 0.9032 - val_loss: 0.0679 - val_accuracy: 0.9225 -
val_recall: 0.9111
Epoch 28/100
477/477 [=====] - 229s 474ms/step - loss: 0.0711 -
accuracy: 0.9141 - recall: 0.9040 - val_loss: 0.0725 - val_accuracy: 0.9219 -
val_recall: 0.9129
Epoch 29/100
477/477 [=====] - 230s 476ms/step - loss: 0.0703 -
accuracy: 0.9202 - recall: 0.9104 - val_loss: 0.0867 - val_accuracy: 0.9159 -
val_recall: 0.9038
Epoch 30/100
477/477 [=====] - 230s 476ms/step - loss: 0.0699 -
accuracy: 0.9209 - recall: 0.9100 - val_loss: 0.0746 - val_accuracy: 0.9153 -
val_recall: 0.9081

Epoch 31/100
477/477 [=====] - 231s 477ms/step - loss: 0.0706 -
accuracy: 0.9145 - recall: 0.9066 - val_loss: 0.0732 - val_accuracy: 0.9075 -
val_recall: 0.9002

Epoch 32/100
477/477 [=====] - 230s 477ms/step - loss: 0.0685 -
accuracy: 0.9240 - recall: 0.9144 - val_loss: 0.0667 - val_accuracy: 0.9243 -
val_recall: 0.9141

Epoch 33/100
477/477 [=====] - 230s 475ms/step - loss: 0.0609 -
accuracy: 0.9270 - recall: 0.9191 - val_loss: 0.0662 - val_accuracy: 0.9201 -
val_recall: 0.9093

Epoch 34/100
477/477 [=====] - 230s 477ms/step - loss: 0.0666 -
accuracy: 0.9256 - recall: 0.9175 - val_loss: 0.0766 - val_accuracy: 0.9105 -
val_recall: 0.9056

Epoch 35/100
477/477 [=====] - 230s 476ms/step - loss: 0.0597 -
accuracy: 0.9272 - recall: 0.9189 - val_loss: 0.0747 - val_accuracy: 0.9044 -
val_recall: 0.8996

Epoch 36/100
477/477 [=====] - 231s 478ms/step - loss: 0.0613 -
accuracy: 0.9264 - recall: 0.9184 - val_loss: 0.0938 - val_accuracy: 0.8984 -
val_recall: 0.8924

Epoch 37/100
477/477 [=====] - 230s 475ms/step - loss: 0.0628 -
accuracy: 0.9275 - recall: 0.9193 - val_loss: 0.0603 - val_accuracy: 0.9387 -
val_recall: 0.9339

Epoch 38/100
477/477 [=====] - 229s 475ms/step - loss: 0.0642 -
accuracy: 0.9255 - recall: 0.9167 - val_loss: 0.0956 - val_accuracy: 0.8966 -
val_recall: 0.8900

Epoch 39/100
477/477 [=====] - 229s 474ms/step - loss: 0.0588 -
accuracy: 0.9289 - recall: 0.9208 - val_loss: 0.1127 - val_accuracy: 0.8798 -
val_recall: 0.8726

Epoch 40/100
477/477 [=====] - 230s 475ms/step - loss: 0.0603 -
accuracy: 0.9315 - recall: 0.9230 - val_loss: 0.0766 - val_accuracy: 0.9141 -
val_recall: 0.9062

Epoch 41/100
477/477 [=====] - 228s 472ms/step - loss: 0.0590 -
accuracy: 0.9298 - recall: 0.9207 - val_loss: 0.0647 - val_accuracy: 0.9309 -
val_recall: 0.9249

Epoch 42/100
477/477 [=====] - 230s 475ms/step - loss: 0.0563 -
accuracy: 0.9359 - recall: 0.9288 - val_loss: 0.0704 - val_accuracy: 0.9279 -
val_recall: 0.9135

```

Epoch 43/100
477/477 [=====] - 230s 475ms/step - loss: 0.0578 -
accuracy: 0.9296 - recall: 0.9229 - val_loss: 0.0737 - val_accuracy: 0.9177 -
val_recall: 0.9129
Epoch 44/100
477/477 [=====] - 229s 474ms/step - loss: 0.0542 -
accuracy: 0.9338 - recall: 0.9264 - val_loss: 0.0622 - val_accuracy: 0.9213 -
val_recall: 0.9129
Epoch 45/100
477/477 [=====] - 229s 474ms/step - loss: 0.0566 -
accuracy: 0.9338 - recall: 0.9278 - val_loss: 0.0670 - val_accuracy: 0.9195 -
val_recall: 0.9135
Epoch 46/100
477/477 [=====] - 230s 476ms/step - loss: 0.0570 -
accuracy: 0.9327 - recall: 0.9258 - val_loss: 0.0641 - val_accuracy: 0.9219 -
val_recall: 0.9153
Epoch 47/100
477/477 [=====] - 229s 474ms/step - loss: 0.0529 -
accuracy: 0.9348 - recall: 0.9292 - val_loss: 0.0921 - val_accuracy: 0.9020 -
val_recall: 0.8960

```

```

[13]: transfer_model.save("./
      ↳tranferlearning_densenet169_with_balanced_focal_loss_3_class.h5")

```

```


[14]: def plot_metrics(history,metrics=[]): #retorna una lista de tuplas
      fig,axes=plt.subplots(1,len(metrics))
      fig.set_size_inches(15,5)
      graph=pd.DataFrame(history)
      for i,ax in enumerate(axes.flat):
          graph[list(metrics[i])].plot(kind="line",style="--",ax=ax)
          ax.set_title(" VS ".join(list(metrics[i])))
          ax.grid(True)
      plt.show()

```

```

[15]: metrics=[("accuracy","val_accuracy"),("recall","val_recall"),("loss","val_loss")]
      plot_metrics(history_model.history,metrics=metrics)

```



```
[16]: #obtenemos el numero de epocas donde se detuvo y lo configuramos como un epoch
      ↪ inicial para el siguiente
      #entrenamiento
      EPOCH_STOP=len(history_model.epoch)
      print("El modelo se entreno en",EPOCH_STOP,"Epochs")
```

El modelo se entreno en 47 Epochs

- El entrenamiento del modelo se detuvo en 47 epochs lo que nos dice que la funcion de perdida en la data de validacion no mejoro por 10 epochs consecutivos, probablemente, ya no mejore para futuras epocas.

Ahora veamos el rendimiento del modelo base en los datos de entrenamiento y validacion

```
[17]: transfer_model.evaluate(train_generator)

477/477 [=====] - 244s 512ms/step - loss: 0.0458 -
accuracy: 0.9466 - recall: 0.9389
```

```
[17]: [0.04583195969462395, 0.9465808868408203, 0.938902735710144]
```

```
[18]: transfer_model.evaluate(validation_generator)

53/53 [=====] - 8s 146ms/step - loss: 0.0606 -
accuracy: 0.9374 - recall: 0.9327
```

```
[18]: [0.0605820007622242, 0.9374262094497681, 0.9327036738395691]
```

- Mostramos la matriz de confusion en los datos de entrenamiento y validacion

Para el conjunto de entrenamiento

```
[19]: y_true=train_generator.classes
      predictions=transfer_model.predict(train_generator)
      y_pred=np.argmax(predictions,axis=1)
```

```
[20]: print("Indices de clase")
      for idx,clase in train_generator.class_indices.items():
          print(idx,":",clase)
```

Indices de clase

COVID : 0

Enfermedades Pulmonares No COVID19 : 1


Normal : 2

```
[21]: #obtenemos la matriz de confusion de sklearn
      from sklearn.metrics import confusion_matrix
      import seaborn as sns; sns.set()
      classes=train_generator.class_indices.keys()
```

```

mat_train=confusion_matrix(y_true,y_pred)
sns.heatmap(mat_train,square=True,annot=True,fmt="0.
    ↪2f",cmap="coolwarm",xticklabels=classes,yticklabels=classes)
plt.xticks(rotation=60)
plt.title("Matriz de confusion en los datos de entrenamiento")
plt.show()

```



- Reporte de clasificacion para el conjunto de entrenamiento

```

[22]: report=classification_report(y_true,y_pred,target_names=list(train_generator.
    ↪class_indices.keys()))
print(report)

```


	precision	recall	f1-score	support
COVID	0.17	0.17	0.17	2604
Enfermedades Pulmonares No COVID19	0.34	0.34	0.34	5296
Normal	0.48	0.49	0.48	7338

accuracy				0.38	15238
macro avg	0.33	0.33	0.33	0.33	15238
weighted avg	0.38	0.38	0.38	0.38	15238

- Para el conjunto de validacion

```
[23]: y_true=validation_generator.classes
      predictions=transfer_model.predict(validation_generator)
      y_pred=np.argmax(predictions,axis=1)
```

```
[24]: classes=validation_generator.class_indices.keys()
      mat_val=confusion_matrix(y_true,y_pred)
      sns.heatmap(mat_val,square=True,annot=True,fmt="0.
      ↪2f",cmap="coolwarm",xticklabels=classes,yticklabels=classes)
      plt.xticks(rotation=60)
      plt.title("Matriz de confusion en los datos de Validacion")
      plt.show()
```



- Reporte de clasificacion para los datos de validacion

```
[25]: report=classification_report(y_true,y_pred,target_names=list(train_generator.
      ↪class_indices.keys()))
      print(report)
```

	precision	recall	f1-score	support
COVID	0.16	0.16	0.16	289
Enfermedades Pulmonares No COVID19	0.35	0.35	0.35	589
Normal	0.47	0.47	0.47	816
accuracy			0.37	1694
macro avg	0.32	0.32	0.32	1694
weighted avg	0.37	0.37	0.37	1694

1.0.1 RESULTADOS FINALES: MODELO TRANSFER LEARNING CON PERDIDA FOCAL PARA EL BALANCEO DE CLASES. USANDO 3 CLASES

- El modelo base ha obtenido un puntaje de accuracy **ACC=94.66%** y recall **RECALL=93.89%** en el **conjunto de entrenamiento**.
- El modelo ha obtendio un puntaje de accuracy **ACC=93.74%** y recall **RECALL=93.27%** en el **conjunto de validacion**.

IMPORTANTE: El modelo ha alcanzado equilibrio entre los datos de entrenamiento y validacion, lo que significa que es un modelo final con aproximadamente 93% de precision. Este modelo pasará a la **Fase de evaluación del modelo**.